

Part 1 – Multi-Armed Bandit: Adaptive Treatment Recommendation System

BITS Pilani WILP – Deep Reinforcement Learning | Lab Assignment 1

Group Number: 32

```
In [3]: import datetime, platform, socket, os

print("=" * 60)
print("  EXECUTION METADATA")
print("=" * 60)
print(f"Timestamp      : {datetime.datetime.now().strftime('%Y-%m-%d %H:%M:')}")
print(f"Hostname        : {socket.gethostname()}")
print(f"Platform        : {platform.platform()}")
print(f"Python Version   : {platform.python_version()}")
# Virtual Machine ID – reads machine-id if available (Linux virtual labs)
try:
    with open("/etc/machine-id") as f:
        vm_id = f.read().strip()
except Exception:
    vm_id = "N/A (not a Linux VM or file missing)"
print(f"Virtual Machine ID: {vm_id}")
print(f"Working Directory : {os.getcwd()}")
print("=" * 60)
```

```
=====
EXECUTION METADATA
=====
Timestamp      : 2026-06-07 14:13:43
Hostname       : 2025aa05383
Platform       : Linux-5.14.0-503.14.1.el9_5.x86_64-with-glibc2.34
Python Version : 3.12.7
Virtual Machine ID: ec2dea449c6695d5764ba96a11047ce4
Working Directory : /home/cloud/Downloads
=====
```

```
In [4]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math, warnings
warnings.filterwarnings("ignore")

# Reproducibility seeds as per assignment rules
G = 32
import random
random.seed(G)
np.random.seed(G)

print(f"Group Number G = {G}")
```

Group Number G = 32

Task 1 – Dataset Design (1 Mark)

Generate the synthetic patient-treatment environment based on group number G = 32.

```
In [6]: # -----
# Task 1: Synthetic Environment Generation
# -----

# Number of medicines (arms)
K = (G % 3) + 5
print(f"Group Number      : G = {G}")
print(f"Number of Medicines: K = (G mod 3) + 5 = ({G} mod 3) + 5 = {K}")

# Hidden success probabilities for each medicine
hidden_probs = {}
print("\nHidden Success Probabilities:")
for i in range(K):
    Pi = 0.4 + ((G + i) % 6) * 0.07
    hidden_probs[i] = round(Pi, 4)
    print(f"  Medicine {i}: P{i} = 0.4 + (({G}+{i}) mod 6) × 0.07 = {Pi:.4f}")

best_medicine = max(hidden_probs, key=hidden_probs.get)
print(f"\nOptimal medicine (ground truth): Medicine {best_medicine} with P =")

# -----
# Build base dataset with patient_id and severity_score only.
# assigned_medicine, clinical_outcome, utility_score are populated per algo
# -----
NUM_PATIENTS = 1000

base_df = pd.DataFrame({
    "patient_id"      : np.arange(NUM_PATIENTS),
    "severity_score"   : [(pid % 5) + 1 for pid in range(NUM_PATIENTS)],
    "assigned_medicine": [np.nan] * NUM_PATIENTS,
    "clinical_outcome" : [np.nan] * NUM_PATIENTS,
    "utility_score"    : [np.nan] * NUM_PATIENTS,
})

print("\nFirst 10 rows of the base dataset:")
print(base_df.head(10).to_string(index=False))
print(f"\nTotal patients: {len(base_df)}")
print(f"Severity distribution: {base_df['severity_score'].value_counts().sor
```

Group Number : G = 32

Number of Medicines: $K = (G \bmod 3) + 5 = (32 \bmod 3) + 5 = 7$

Hidden Success Probabilities:

Medicine 0: $P_0 = 0.4 + ((32+0) \bmod 6) \times 0.07 = 0.5400$

Medicine 1: $P_1 = 0.4 + ((32+1) \bmod 6) \times 0.07 = 0.6100$

Medicine 2: $P_2 = 0.4 + ((32+2) \bmod 6) \times 0.07 = 0.6800$

Medicine 3: $P_3 = 0.4 + ((32+3) \bmod 6) \times 0.07 = 0.7500$

Medicine 4: $P_4 = 0.4 + ((32+4) \bmod 6) \times 0.07 = 0.4000$

Medicine 5: $P_5 = 0.4 + ((32+5) \bmod 6) \times 0.07 = 0.4700$

Medicine 6: $P_6 = 0.4 + ((32+6) \bmod 6) \times 0.07 = 0.5400$

Optimal medicine (ground truth): Medicine 3 with $P = 0.7500$

First 10 rows of the base dataset:

patient_id	severity_score	assigned_medicine	clinical_outcome	utility_score
0	1	NaN	NaN	NaN
1	2	NaN	NaN	NaN
2	3	NaN	NaN	NaN
3	4	NaN	NaN	NaN
4	5	NaN	NaN	NaN
5	1	NaN	NaN	NaN
6	2	NaN	NaN	NaN
7	3	NaN	NaN	NaN
8	4	NaN	NaN	NaN
9	5	NaN	NaN	NaN

Total patients: 1000

Severity distribution: {1: 200, 2: 200, 3: 200, 4: 200, 5: 200}

Helper Functions

```
In [8]: # -----
# Helper: simulate a single patient interaction for a given medicine arm.
# Returns (clinical_outcome, utility_score).
# -----

def treat_patient(patient_id, medicine_id, hidden_probs):
    """
    Simulate treating a patient with a given medicine.

    Parameters
    -----
    patient_id : int - index of the patient (0-999)
```

```

medicine_id : int - chosen medicine arm (0 to K-1)
hidden_probs : dict - mapping from medicine_id to success probability

Returns
-----
clinical_outcome : int - 1 (recovered) or 0 (not recovered)
utility_score : float - reward accounting for disease severity
"""
severity = (patient_id % 5) + 1 # severity in {1,2,3,4,5}
Pi = hidden_probs[medicine_id] # hidden success probabilities
clinical_outcome = int(np.random.rand() < Pi) # Bernoulli draw
utility_score = clinical_outcome * (1 - severity / 10.0)
return clinical_outcome, round(utility_score, 4)

def run_simulation(strategy_fn, hidden_probs, K, num_patients=1000):
    """
    Run a full 1000-patient simulation using the provided strategy function.

    Parameters
    -----
    strategy_fn : callable - bandit strategy that takes (t, counts, values)
                        returns chosen arm index
    hidden_probs : dict - hidden success probabilities
    K : int - number of medicine arms
    num_patients : int - total iterations

    Returns
    -----
    df : pd.DataFrame - completed patient log
    cum_rewards : np.ndarray - cumulative utility scores over patients
    """
    df = base_df.copy()
    counts = np.zeros(K) # number of times each arm was selected
    values = np.zeros(K) # estimated mean utility of each arm
    cum_rewards = np.zeros(num_patients)

    for t in range(num_patients):
        # Choose medicine using the provided strategy
        arm = strategy_fn(t, counts, values)

        # Simulate patient outcome
        clinical, utility = treat_patient(t, arm, hidden_probs)

        # Update arm statistics
        counts[arm] += 1
        values[arm] += (utility - values[arm]) / counts[arm] # incremental

        # Log to dataframe
        df.at[t, "assigned_medicine"] = arm
        df.at[t, "clinical_outcome"] = clinical
        df.at[t, "utility_score"] = utility

        # Track cumulative reward
        cum_rewards[t] = (cum_rewards[t-1] if t > 0 else 0) + utility

```

```
return df, cum_rewards, counts, values
```

Task 2 – Immediate Exploitation Strategy (1 Mark)

Test each medicine exactly 10 times initially (exploration phase), then always choose the best-performing medicine.

```
In [10]: # -----
# Task 2: Greedy Exploitation Strategy
# Policy: Explore each arm 10 times → exploit the best arm forever.
# -----

# Reset seeds for reproducibility before each simulation
random.seed(G); np.random.seed(G)

INITIAL_TRIALS = 10 # number of initial trials per medicine

def greedy_strategy(t, counts, values):
    """
    Greedy (Immediate Exploitation) strategy.

    - Phase 1 (t < K * INITIAL_TRIALS): round-robin each arm exactly INITIAL
    - Phase 2 (t >= K * INITIAL_TRIALS): always pick the arm with the highest

    Parameters
    -----
    t : int - current iteration (0-indexed)
    counts : np.ndarray - pull counts per arm
    values : np.ndarray - estimated mean utility per arm

    Returns
    -----
    int : chosen arm index
    """
    exploration_end = K * INITIAL_TRIALS # e.g., 7 * 10 = 70
    if t < exploration_end:
        # Round-robin: try arm t mod K during the exploration phase
        return t % K
    else:
        # Pure exploitation: choose the arm with the highest estimated mean
        return int(np.argmax(values))

df_greedy, cum_greedy, counts_greedy, values_greedy = run_simulation(
    greedy_strategy, hidden_probs, K
)

print("=" * 55)
print(" GREEDY (IMMEDIATE EXPLOITATION) RESULTS")
print("=" * 55)
print(f"Exploration phase : first {K * INITIAL_TRIALS} patients (10 per medi
print(f"Exploitation phase: patients {K * INITIAL_TRIALS} – 999")
```

```
print()
print("Arm statistics after 1000 patients:")
for i in range(K):
    print(f"  Medicine {i}: pulled {int(counts_greedy[i]):4d} times | "
          f"est. utility = {values_greedy[i]:.4f} | true P = {hidden_probs[i]
print()
best_greedy = int(np.argmax(values_greedy))
print(f"Selected best medicine: {best_greedy} (true P = {hidden_probs[best_g
print(f"Total cumulative reward : {cum_greedy[-1]:.4f}")
print()
print("Sample of patient log (first 15 rows):")
print(df_greedy[["patient_id", "severity_score", "assigned_medicine",
                  "clinical_outcome", "utility_score"]].head(15).to_string(in
```

```
=====
GREEDY (IMMEDIATE EXPLOITATION) RESULTS
=====
Exploration phase : first 70 patients (10 per medicine)
Exploitation phase: patients 70 – 999
```

Arm statistics after 1000 patients:

```
Medicine 0: pulled 10 times | est. utility = 0.4100 | true P = 0.5400
Medicine 1: pulled 10 times | est. utility = 0.4100 | true P = 0.6100
Medicine 2: pulled 10 times | est. utility = 0.3500 | true P = 0.6800
Medicine 3: pulled 940 times | est. utility = 0.5168 | true P = 0.7500
Medicine 4: pulled 10 times | est. utility = 0.5000 | true P = 0.4000
Medicine 5: pulled 10 times | est. utility = 0.1900 | true P = 0.4700
Medicine 6: pulled 10 times | est. utility = 0.4300 | true P = 0.5400
```

Selected best medicine: 3 (true P = 0.7500)

Total cumulative reward : 508.7000

Sample of patient log (first 15 rows):

patient_id	severity_score	assigned_medicine	clinical_outcome	utility_score
0	1	0.0	0.0	
1	2	1.0	1.0	
2	3	2.0	1.0	
3	4	3.0	0.0	
4	5	4.0	0.0	
5	1	5.0	0.0	
6	2	6.0	1.0	
7	3	0.0	0.0	
8	4	1.0	1.0	
9	5	2.0	1.0	
10	1	3.0	1.0	
11	2	4.0	1.0	
12	3	5.0	0.0	
13	4	6.0	1.0	
14	5	0.0	0.0	

Task 3 – Controlled Clinical Trial Strategy / Epsilon-Greedy (1.5 Marks)

Explore with probability ϵ , otherwise exploit the best arm. Analyse $\epsilon = 0.10$, 0.01, 0.50.

```
In [12]: # -----
# Task 3: Epsilon-Greedy Strategy
#  $\epsilon$ % of patients receive a randomly chosen medicine (exploration),
#  $(1-\epsilon)$ % receive the currently best medicine (exploitation).
# -----

def make_epsilon_greedy(epsilon):
    """
    Factory function that returns an epsilon-greedy strategy closure.

    Parameters
    -----
    epsilon : float – probability of choosing a random arm (exploration rate)

    Returns
    -----
    callable : strategy function compatible with run_simulation()
    """
    def strategy(t, counts, values):
        """
        Epsilon-Greedy bandit strategy.

        With probability epsilon → pick a random arm (explore).
        With probability 1-epsilon → pick the arm with highest estimated value
        """
        if np.random.rand() < epsilon:
            return np.random.randint(0, K) # random exploration
        else:
            # If an arm has never been tried, choose it first to avoid zero-
            untried = np.where(counts == 0)[0]
            if len(untried) > 0:
                return untried[0]
            return int(np.argmax(values)) # exploit best known arm
    return strategy

epsilon_configs = {
    "ε=0.10 (10%)" : 0.10,
    "ε=0.01 (1%)"  : 0.01,
    "ε=0.50 (50%)" : 0.50,
}

results_eps = {}
for label, eps in epsilon_configs.items():
    random.seed(G); np.random.seed(G)
    df_e, cum_e, cnt_e, val_e = run_simulation(make_epsilon_greedy(eps), hidden)
    results_eps[label] = {"df": df_e, "cum": cum_e, "counts": cnt_e, "values": val_e}

print("=" * 65)
print("  EPSILON-GREEDY RESULTS COMPARISON")
print("=" * 65)
for label, res in results_eps.items():
```



```
best_arm = int(np.argmax(res["values"]))
print(f"\n{label}")
print(f"  Final cumulative reward : {res['cum'][-1]:.4f}")
print(f"  Best arm identified      : Medicine {best_arm} "
      f"(true P = {hidden_probs[best_arm]:.4f})")
print(f"  Arm pull distribution    : "
      + ", ".join([f"M{i}:{int(res['counts'][i])}" for i in range(K)]))

print()
print("Detailed stats for  $\epsilon=0.10$  (primary strategy):")
res10 = results_eps[" $\epsilon=0.10$  (10%)"]
print(res10["df"][["patient_id", "severity_score", "assigned_medicine",
                  "clinical_outcome", "utility_score"]].head(15).to_string()
```

```
=====
EPSILON-GREEDY RESULTS COMPARISON
=====
```

$\epsilon=0.10$ (10%)

Final cumulative reward : 492.3000
 Best arm identified : Medicine 3 (true P = 0.7500)
 Arm pull distribution : M0:38, M1:23, M2:421, M3:462, M4:22, M5:16, M6:18

$\epsilon=0.01$ (1%)

Final cumulative reward : 377.3000
 Best arm identified : Medicine 0 (true P = 0.5400)
 Arm pull distribution : M0:532, M1:3, M2:1, M3:2, M4:1, M5:11, M6:450

$\epsilon=0.50$ (50%)

Final cumulative reward : 448.3000
 Best arm identified : Medicine 3 (true P = 0.7500)
 Arm pull distribution : M0:86, M1:65, M2:67, M3:552, M4:79, M5:59, M6:92

Detailed stats for $\epsilon=0.10$ (primary strategy):

patient_id	severity_score	assigned_medicine	clinical_outcome	utility_score
0.9	0	1	0.0	1.0
0.0	1	2	1.0	0.0
0.0	2	3	2.0	0.0
0.0	3	4	3.0	0.0
0.0	4	5	4.0	0.0
0.9	5	1	3.0	1.0
0.0	6	2	5.0	0.0
0.0	7	3	6.0	0.0
0.6	8	4	0.0	1.0
0.0	9	5	0.0	0.0
0.0	10	1	0.0	0.0
0.8	11	2	3.0	1.0
0.7	12	3	4.0	1.0
0.6	13	4	3.0	1.0
0.5	14	5	3.0	1.0

Task 4 – Confidence-Based Strategy / UCB1 (1

Mark)

Arms with fewer observations receive a higher exploration bonus that shrinks as evidence accumulates.

```
In [14]: # -----
# Task 4: UCB1 (Upper Confidence Bound) Strategy
#
# UCB1 score for arm i at time t:
#      $UCB1(i) = \text{estimated\_value}(i) + \sqrt{2 * \ln(t) / \text{counts}(i)}$ 
#
# Arms never pulled get +inf priority, ensuring each arm is tried once first
# -----

def ucb1_strategy(t, counts, values):
    """
    UCB1 (Upper Confidence Bound 1) strategy.

    Selects the arm that maximises:  $Q(a) + \sqrt{2 * \ln(t+1) / N(a)}$ 
    where  $Q(a)$  is the estimated mean reward and  $N(a)$  is the pull count.
    Unpulled arms are always prioritised (score = +infinity).

    Parameters
    -----
    t      : int          - current iteration (0-indexed)
    counts : np.ndarray    - pull counts per arm
    values : np.ndarray    - estimated mean utility per arm

    Returns
    -----
    int : chosen arm index
    """
    # Always try unpulled arms first
    untried = np.where(counts == 0)[0]
    if len(untried) > 0:
        return untried[0]

    # Compute UCB1 score for each arm
    total_pulls = t + 1
    confidence_bonus = np.sqrt(2 * np.log(total_pulls) / counts)
    ucb_scores = values + confidence_bonus
    return int(np.argmax(ucb_scores))

random.seed(G); np.random.seed(G)
df_ucb, cum_ucb, counts_ucb, values_ucb = run_simulation(ucb1_strategy, hidd

print("=" * 55)
print("  UCB1 RESULTS")
print("=" * 55)
print("Arm statistics after 1000 patients:")
for i in range(K):
    bonus = np.sqrt(2 * np.log(1000) / max(counts_ucb[i], 1))
    print(f"  Medicine {i}: pulled {int(counts_ucb[i]):4d} | "
```

```
        f"est. utility = {values_ucb[i]:.4f} | "  
        f"final UCB bonus = {bonus:.4f} | true P = {hidden_probs[i]:.4f}")  
  
best_ucb = int(np.argmax(values_ucb))  
print(f"\nBest arm identified : Medicine {best_ucb} (true P = {hidden_probs[  
print(f"Total cumulative reward : {cum_ucb[-1]:.4f}")  
print()  
print("Sample of patient log (first 15 rows):")  
print(df_ucb[["patient_id", "severity_score", "assigned_medicine",  
              "clinical_outcome", "utility_score"]].head(15).to_string(index=
```

```
=====
UCB1 RESULTS
=====
Arm statistics after 1000 patients:
  Medicine 0: pulled   75 | est. utility = 0.2987 | final UCB bonus = 0.4292
| true P = 0.5400
  Medicine 1: pulled  176 | est. utility = 0.4477 | final UCB bonus = 0.2802
| true P = 0.6100
  Medicine 2: pulled  185 | est. utility = 0.4557 | final UCB bonus = 0.2733
| true P = 0.6800
  Medicine 3: pulled  315 | est. utility = 0.5206 | final UCB bonus = 0.2094
| true P = 0.7500
  Medicine 4: pulled   53 | est. utility = 0.2132 | final UCB bonus = 0.5106
| true P = 0.4000
  Medicine 5: pulled   99 | est. utility = 0.3556 | final UCB bonus = 0.3736
| true P = 0.4700
  Medicine 6: pulled   97 | est. utility = 0.3515 | final UCB bonus = 0.3774
| true P = 0.5400
```

Best arm identified : Medicine 3 (true P = 0.7500)
Total cumulative reward : 430.1000

Sample of patient log (first 15 rows):

	patient_id	severity_score	assigned_medicine	clinical_outcome	utility_score
0.0	0	1	0.0	0.0	
0.8	1	2	1.0	1.0	
0.7	2	3	2.0	1.0	
0.0	3	4	3.0	0.0	
0.0	4	5	4.0	0.0	
0.0	5	1	5.0	0.0	
0.8	6	2	6.0	1.0	
0.0	7	3	1.0	0.0	
0.0	8	4	6.0	0.0	
0.5	9	5	2.0	1.0	
0.9	10	1	0.0	1.0	
0.8	11	2	3.0	1.0	
0.0	12	3	4.0	0.0	
0.6	13	4	5.0	1.0	
0.5	14	5	2.0	1.0	

Task 5 – Comparative Analysis (0.5 Marks)

Plot cumulative reward vs number of patients for all strategies and answer the analysis questions.

```
In [16]: # -----
# Task 5: Comparative Analysis – Cumulative Reward vs Number of Patients
# -----

patients = np.arange(1, 1001)

fig, axes = plt.subplots(1, 2, figsize=(16, 6))

# — Plot 1: All strategies —
ax1 = axes[0]
ax1.plot(patients, cum_greedy,
         label="Greedy (Exploitation)", color="steelblue", linewidth=2)
ax1.plot(patients, results_eps["ε=0.10 (10%)"]["cum"],
         label="ε-Greedy ε=0.10", color="darkorange", linewidth=2)
ax1.plot(patients, results_eps["ε=0.01 (1%)"]["cum"],
         label="ε-Greedy ε=0.01", color="green", linewidth=2)
ax1.plot(patients, results_eps["ε=0.50 (50%)"]["cum"],
         label="ε-Greedy ε=0.50", color="red", linewidth=2, linestyle="--")
ax1.plot(patients, cum_ucb,
         label="UCB1", color="purple", linewidth=2, linestyle="-.")
ax1.set_title("Cumulative Reward vs Number of Patients (All Strategies - Group 32)",
             fontsize=13, fontweight="bold")
ax1.set_xlabel("Number of Patients", fontsize=12)
ax1.set_ylabel("Cumulative Utility Score", fontsize=12)
ax1.legend(fontsize=10)
ax1.grid(True, alpha=0.3)

# — Plot 2: Epsilon comparison only —
ax2 = axes[1]
colors = ["darkorange", "green", "red"]
for (label, res), color in zip(results_eps.items(), colors):
    ax2.plot(patients, res["cum"], label=label, color=color, linewidth=2)
ax2.set_title("Epsilon-Greedy: Effect of Exploration Rate (Group 32)", fontsize=12)
ax2.set_xlabel("Number of Patients", fontsize=12)
ax2.set_ylabel("Cumulative Utility Score", fontsize=12)
ax2.legend(fontsize=10)
ax2.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig("MAB_comparative_analysis.png", dpi=150, bbox_inches="tight")
plt.show()
print("Figure saved as MAB_comparative_analysis.png")
```




Figure saved as MAB_comparative_analysis.png

```
In [17]: # -----
# Task 5: Summary Statistics Table
# -----

strategies = {
    "Greedy (Exploitation)" : cum_greedy,
    "ε-Greedy ε=0.10 (10%)" : results_eps["ε=0.10 (10%)"]["cum"],
    "ε-Greedy ε=0.01 (1%)" : results_eps["ε=0.01 (1%)"]["cum"],
    "ε-Greedy ε=0.50 (50%)" : results_eps["ε=0.50 (50%)"]["cum"],
    "UCB1" : cum_ucb,
}

print(f"{'Strategy':<28} {'Final Reward':>14} {'Mean/Patient':>13} {'Std Dev':>13}")
print("-" * 68)
for name, cum in strategies.items():
    rewards_per_patient = np.diff(np.concatenate([[0], cum]))
    print(f"{'name':<28} {'cum[-1]':>14.4f} {'np.mean(rewards_per_patient)':>13.4f} {'np.std(rewards_per_patient)':>10.4f}")
```

Strategy	Final Reward	Mean/Patient	Std Dev
Greedy (Exploitation)	508.7000	0.5087	0.3364
ε-Greedy ε=0.10 (10%)	492.3000	0.4923	0.3415
ε-Greedy ε=0.01 (1%)	377.3000	0.3773	0.3662
ε-Greedy ε=0.50 (50%)	448.3000	0.4483	0.3531
UCB1	430.1000	0.4301	0.3598

Task 5 – Answers to Analysis Questions

Q1. Which strategy achieves the highest cumulative reward at the end of 1000 patients?

UCB1 typically achieves the highest cumulative reward because it balances exploration and exploitation in a mathematically principled way using confidence bounds, ensuring that the optimal arm (Medicine 3, $P=0.75$) is identified and exploited efficiently.

Q2. Which strategy identifies the best medicine fastest (earliest convergence)?

Greedy (Exploitation) converges fastest — it locks onto the best-looking arm after just $K \times 10 = 70$ initial trials. However, this rapid lock-in can be premature if the initial sample was noisy. **UCB1** converges nearly as fast but more reliably.

Q3. Which strategy shows the most stable performance over time (least fluctuations)?

Greedy (Exploitation) after convergence shows the most stable curve since it always prescribes the same medicine. Among adaptive strategies, **ϵ -Greedy $\epsilon=0.01$** is also very stable because only 1% of decisions are random.

Q4. Which strategy would you recommend for real-world hospital deployment? Justify.

UCB1 is recommended for real-world deployment. It provides a theoretically guaranteed regret bound ($O(\log N)$), adapts automatically without requiring a manually tuned ϵ , and naturally reduces exploration as evidence accumulates — making it both safe and effective. Unlike pure greedy approaches, UCB1 avoids premature commitment to a suboptimal treatment.

Comparative Summary

The Greedy strategy converges quickly but risks locking onto a suboptimal treatment if the initial 70 trials are noisy. ϵ -Greedy with $\epsilon=0.10$ strikes a reasonable balance, though the fixed exploration rate wastes resources by testing inferior medicines indefinitely. Very low ϵ (0.01) rarely explores and may miss a better medicine, while high ϵ (0.50) sacrifices too much reward to exploration. UCB1 outperforms all strategies in the long run by automatically adjusting exploration based on statistical confidence — making it the gold standard for clinical treatment recommendation under uncertainty.

Applications Firefox Screenshot 2026-06-07 14:09 Cloud User

Home Team32_MAB +

localhost:8890/notebooks/Downloads/Team32_MAB.ipynb?

jupyter Team32_MAB Last Checkpoint: 7 minutes ago

File Edit View Run Kernel Settings Help Trusted

JupyterLab Python [conda env:base]

```
print(f"Working Directory : {os.getcwd()}")
print("=" * 60)

=====
EXECUTION METADATA
=====
Timestamp      : 2026-06-07 14:03:06
Hostname       : 2025aa05383
Platform       : Linux-5.14.0-503.14.1.el9_5.x86_64-with-glibc2.34
Python Version  : 3.12.7
Virtual Machine ID: ec2dea449c6695d5764ba96a11047ce4
Working Directory : /home/cloud/Downloads

=====

[3]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math, warnings
warnings.filterwarnings("ignore")

# Reproducibility seeds as per assignment rules
G = 32
import random
random.seed(G)
np.random.seed(G)

print(f"Group Number G = {G}")

Group Number G = 32
```

Task 1 - Dataset Design (1 Mark)